


Amortization: SBI pays a one-time training cost, then near-free inference

Cumulative inference cost



Per-dataset speed-up $\approx 22,123\times$

